

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
from sympy import *
```

13.3

Consider the polynomial $p(x) = (x - 2)^9$. Plot this polynomial in it's expanded form and in the factored form.

```
In [ ]: def p1(x):
return x ** 9 - 18 * x ** 8 + 144 * x ** 7 - 672 * x ** 6 + 2016 * x ** 5 - 403
```

```
In [ ]: x = np.round(np.linspace(-1.920, 2.080, 4000), decimals=3)
```

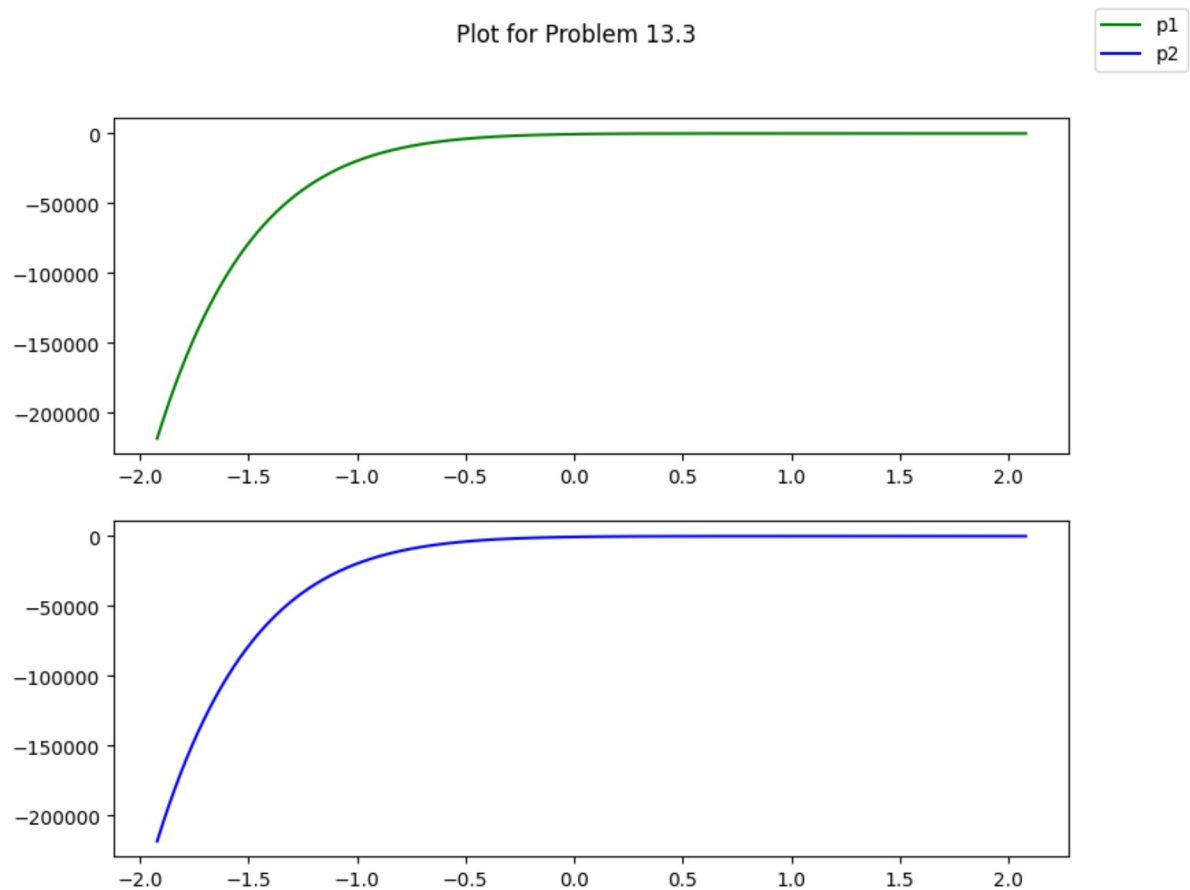
```
In [ ]: x
```

```
Out[ ]: array([-1.92 , -1.919, -1.918, ...,  2.078,  2.079,  2.08 ])
```

```
In [ ]: def p2(x):
return (x - 2) ** 9
```

```
In [ ]: fig, axs = plt.subplots(2, figsize=(9, 7))
axs[0].plot(x, p1(x), color='green', label='p1')
axs[1].plot(x, p2(x), color='blue', label='p2')
fig.suptitle('Plot for Problem 13.3')
fig.legend()
```

```
Out[ ]: <matplotlib.legend.Legend at 0x7fb18b919a60>
```



13.4

The polynomial $p(x) = x^5 - 2x^4 - 3x^3 + 3x^2 - 2x - 1$ has three real zeros. Applying Newton's method to p with initial guess $x_0 = 0$ produces a series of estimates x_1, x_2, x_3, \dots that converge rapidly to a zero $x_* \approx -0.315$.

(a) Compute x_1, \dots, x_6 in floating point arithmetic with $\epsilon_{\text{machine}} \approx 10^{-16}$. How many digits do you estimate are correct in each of these numbers?

```
In [ ]: # Machine epsilon for Numpy's float64 is approximately 10^{-16}
np.finfo(np.float64).eps
```

```
Out[ ]: 2.22044604925031 · 10^{-16}
```

```
In [ ]: def newton(f, dfdx, x0, epsilon=1e-10, max_iters=500):
        # Quick iterative implementation of Newton's method
        iters = 0
        estimates = [x0]
        x_previous = x0
        while np.abs(f(x_previous)) > epsilon and iters < max_iters:
            x_next = x_previous - f(x_previous)/dfdx(x_previous)
            estimates.append(x_next)
            x_previous = x_next
            iters += 1
        return x_next, estimates
```

```
In [ ]: f = lambda x: (x ** 5) - (2 * x ** 4) - (3 * x ** 3) + 3 * x ** 2 - 2 * x - 1
        dfdx = lambda x: (5 * x ** 4) - (8 * x ** 3) - (9 * x ** 2) + (6 * x) - 2
        x0 = np.float64(0.0)
```

```
In [ ]: x_final, floating_point_xs = newton(f, dfdx, x0)
```

Below are the sequence of estimates produced by Newton's method. I would estimate that no more than 15 digits after the decimal are correct (but probably fewer due to the division in the algorithm).

```
In [ ]: floating_point_xs
```

```
Out[ ]: [0.0, - 0.5, - 0.336842105263158, - 0.315728448396289, - 0.315301162703277, - 0.315301162703277]
```

```
In [ ]:
```

(b) Compute x_1, \dots, x_6 again *exactly* with the aid of a symbolic algebra system such as Maple or Mathematica. Each x_j is a rational number. How many digits are there in the numerator and the denominator for each j ?

```
In [ ]: # Using Sympy
        x= symbols('x')
        init_printing(use_unicode=True)
```

```
In [ ]: f = (x ** 5) - (2 * x ** 4) - (3 * x ** 3) + 3 * x ** 2 - 2 * x - 1
        dfdx = diff(f, x)
```

```
In [ ]: dfdx
```

```
Out[ ]: 5x4 - 8x3 - 9x2 + 6x - 2
```

```
In [ ]: def sym_newton(f, x0, epsilon=1e-10, max_iters=500):
# symx_previous, symx_next = symbols('symx_previous symxnext dfdx')
dfdx = diff(f)
iters = 0
estimates = [x0]
fs = list()
dfdxs = list()
x_previous = x0
symx_previous = x
while np.abs(f.subs(x, x_previous)) > epsilon and iters < max_iters:
    x_next = x_previous - f.subs(x, x_previous)/dfdx.subs(x, x_previous)
    fs.append(f.subs(x, x_previous))
    dfdxs.append(dfdx.subs(x, x_previous))
    estimates.append(x_next)
    x_previous = x_next
    iters += 1
return x_next, estimates, fs, dfdxs
```

```
In [ ]: result, symbolic_xs, fs, dfdxs = sym_newton(f, x0=0.0)
```

```
In [ ]: print('Floating point estimates: \n' + str(floating_point_xs))
print('Symbolic estimates \n' + str(symbolic_xs))
```

Floating point estimates:

```
[0.0, -0.5, -0.3368421052631579, -0.31572844839628944, -0.31530116270327685, -0.31530098645936266]
```

Symbolic estimates

```
[0.0, -0.5000000000000000, -0.336842105263158, -0.315728448396289, -0.315301162703277, -0.315300986459363]
```

```
In [ ]: fs
```

```
Out[ ]: [-1.0, 0.96875, 0.0986450190238758, 0.00191853843116641, 7.90693465868575 · 10-7]
```

```
In [ ]: dfdxs
```

```
Out[ ]: [-2, - 5.9375, - 4.67209539521643, - 4.49006007582294, - 4.4863589749122]
```

Taking a look at the values computed for f and df/dx at each step j , it appears that there are at least 15 digits after the second iteration.